

SparkMLlib

实验目的： 熟练掌握sparkmllib的使用

准备工作：

```
144 # pyspark
145 export PYSARK_DRIVER_PYTHON=jupyter
146 # export PYSARK_DRIVER_PYTHON=/usr/bin/python3
147 export PYSARK_DRIVER_PYTHON_OPTS='notebook'
148 export PYSARK_PYTHON=python3
149 # spark streaming
150 export PYTHONPATH=$SPARK_HOME/python:$SPARK_HOME/python/lib/py4j-0.10.7-src.zip:$PYTHONPATH
```

激活环境变量

```
chen@ubuntu:~$ . .bashrc
```

启动hadoop

```
chen@ubuntu:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as chen in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
```

启动spark

```
chen@ubuntu:~$ /apps/spark/sbin/start-all.sh
starting org.apache.spark.deploy.master.Master, logging to /apps/spark/logs/spark-chen-org.apache.spark.deploy.master.Master-1-ubuntu.out
```

检查

```
chen@ubuntu:~$ jps
12832 DataNode
14977 Worker
12658 NameNode
15028 Jps
14839 Master
13640 NodeManager
13354 ResourceManager
13119 SecondaryNameNode
```

上传数据

```
hadoop fs -mkdir /input/mllib
hadoop fs -put /data/train.tsv /input/mllib/
hadoop fs -put /data/test.tsv /input/mllib/
```

```
chen@ubuntu:/data$ hadoop fs -mkdir /input/mllib
chen@ubuntu:/data$ hadoop fs -put /data/train.tsv /input/mllib/
chen@ubuntu:/data$ hadoop fs -put /data/test.tsv /input/mllib/
```

检查

```
hadoop fs -ls /input/mllib
```

```
chen@ubuntu:/data$ hadoop fs -ls /input/mllib
Found 2 items
-rw-r--r-- 1 chen supergroup 9428650 2020-12-15 19:30 /input/mllib/test.tsv
-rw-r--r-- 1 chen supergroup 21972916 2020-12-15 19:29 /input/mllib/train.tsv
chen@ubuntu:/data$
```

启动pyspark

创建目录~/pyspark-workspace/mllib/，在该目录中启动 Pyspark，创建一个 Notebook。

```
chen@ubuntu:~$ cd -  
/home/chen/pyspark-workspace/mllib  
chen@ubuntu:~/pyspark-workspace/mllib$ pyspark  
[I 19:33:49.214 NotebookApp] Writing notebook server cookie secret to /run/user/1000/jupyter/notebook_cookie_secret  
[I 19:33:49.547 NotebookApp] Serving notebooks from local directory: /home/chen/pyspark-workspace/mllib  
[I 19:33:49.547 NotebookApp] 0 active kernels  
[I 19:33:49.547 NotebookApp] The Jupyter Notebook is running at:  
[I 19:33:49.548 NotebookApp] http://10.0.0.135:8888/  
[I 19:33:49.548 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
```

读取数据

读取hdfs上的train.tsv文件，并查看数据项

```
In [1]: from pyspark import SparkContext, SparkConf
        from pyspark.sql import SQLContext
        conf = SparkConf()
        conf.setAppName('Streaming').set('spark.io.compression.codec', 'snappy')
        conf.setMaster('local[2]')
        sc = SparkContext(conf = conf)
        sqlContext = SQLContext(sc)

In [2]: row_df = sqlContext.read.format("csv")\
        .option("header", "true")\
        .option("delimiter", "\t")\
        .load("/input/mllib/train.tsv")
        print(row_df.count())

7395

In [3]: row_df.printSchema()

root
|-- url: string (nullable = true)
|-- urlid: string (nullable = true)
|-- boilerplate: string (nullable = true)
|-- alchemy_category: string (nullable = true)
|-- alchemy_category_score: string (nullable = true)
|-- avglinksize: string (nullable = true)
|-- commonlinkratio_1: string (nullable = true)
|-- commonlinkratio_2: string (nullable = true)
|-- commonlinkratio_3: string (nullable = true)
|-- commonlinkratio_4: string (nullable = true)
|-- compression_ratio: string (nullable = true)
|-- embed_ratio: string (nullable = true)
|-- framebased: string (nullable = true)
|-- frameTagRatio: string (nullable = true)
|-- hasDomainLink: string (nullable = true)
|-- html_ratio: string (nullable = true)
|-- image_ratio: string (nullable = true)
|-- is_news: string (nullable = true)
|-- lengthyLinkDomain: string (nullable = true)
|-- linkwordscore: string (nullable = true)
|-- news_front_page: string (nullable = true)
|-- non_markup_alphanum_characters: string (nullable = true)
|-- numberOfLinks: string (nullable = true)
|-- numwords_in_url: string (nullable = true)
|-- parametrizedLinkRatio: string (nullable = true)
|-- spelling_errors_ratio: string (nullable = true)
|-- label: string (nullable = true)

In [4]: row_df.show(1)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

url urlid	boilerplate alchemy_category alchemy_category_score avglinks commonlinkratio_1 commonlinkratio_2 commonlinkratio_3 commonlinkratio_4 compression_ratio embed_ratio framebased frameTagRatio hasDomainLink html_ratio image_ratio is_news lengthyLinkDomain linkwordscore news_fron ont_page non_markup_alphanum_characters numberOfLinks numwords_in_url parametrizedLinkRatio spelling_errors_ratio label
http://www.bloomb... 4042	{ ""title"": ""IBM... business 0.789131 2.055555556 0.676470588 0.205882353 0.047058824 0.023529412 0.443783175 0 0 0.09077381 0 0.245831182 0.003883495 1 1 24 0 5424 170 8 0.152941176 0.079129575 0

only showing top 1 row

```
In [5]: row_df.select('url','urlid','boilerplate','alchemy_category','alchemy_category_score','is_news','label').show(10)
```

url urlid	boilerplate alchemy_category alchemy_category_score is_news label
http://www.bloomb... 4042	{ ""title"": ""IBM... business 0.789131 1 0
http://www.popsci... 8471	{ ""title"": ""The... recreation 0.574147 1 1
http://www.menshe... 1164	{ ""title"": ""Fru... health 0.996526 1 1
http://www.dumbli... 6684	{ ""title"": ""10 ... health 0.801248 1 1
http://bleacherre... 9006	{ ""title"": ""The... sports 0.719157 1 0
http://www.conven... 7018	{ ""url"": ""conve... ? ? ? 0
http://gofashionl... 8685	{ ""title"": ""fas...arts_entertainment 0.22111 1 1
http://www.inside... 3402	{ ""url"": ""insid... ? ? ? 0
http://www.valetm... 477	{ ""title"": ""Val... ? ? 1 1
http://www.howswe... 6731	{ ""url"": ""howsw... ? ? ? 1

only showing top 10 rows

编写 DataFrames UDF 用户自定义函数，将数据中的 ? 转换为 0

```
In [6]: from pyspark.sql.functions import udf
def replace_question(x):
    return ("0" if x=="?" else x)
replace_question=udf(replace_question)
```

导入 col 模块及 pyspark.sql.types 模块，后续可以使用 col 模块读取字段数据，使用 pyspark.sql.types 模块转换数据类型

```
In [7]: from pyspark.sql.functions import col
import pyspark.sql.types
```

使用 replace_question UDF 用户自定义函数，将 row_df DataFrame 第 4 个字段至最后一个字段转换为 double。其中，最后一个字段为 label，其余是 feature。

```
In [8]: df = row_df.select(
    ['url','alchemy_category'] +
    [replace_question(col(column)).cast("double").alias(column) for column in row_df.columns[4:]]
)
```

说明: - 用 `row_df.select` 选取字段

- 选取字段 `['url','alchemy_category']`, 不需要转换
- `for column in row_df.columns[4:]` 读取第 4 个字段至最后一个字段
- `col(column)` 读取字段数据并调用 `replace_question` 自定义函数删除问号 “?”
- `.cast("double")` 转换为 `double`
- `.alias(column)` 把别名设置为原来的字段名

查看使用 `replace_question` UDF 转换后的字段

```
In [9]: df.printSchema()
```

```
root
|-- url: string (nullable = true)
|-- alchemy_category: string (nullable = true)
|-- alchemy_category_score: double (nullable = true)
|-- avglinksize: double (nullable = true)
|-- commonlinkratio_1: double (nullable = true)
|-- commonlinkratio_2: double (nullable = true)
|-- commonlinkratio_3: double (nullable = true)
|-- commonlinkratio_4: double (nullable = true)
|-- compression_ratio: double (nullable = true)
|-- embed_ratio: double (nullable = true)
|-- framebased: double (nullable = true)
|-- frameTagRatio: double (nullable = true)
|-- hasDomainLink: double (nullable = true)
|-- html_ratio: double (nullable = true)
|-- image_ratio: double (nullable = true)
|-- is_news: double (nullable = true)
|-- lengthyLinkDomain: double (nullable = true)
|-- linkwordscore: double (nullable = true)
|-- news_front_page: double (nullable = true)
|-- non_markup_alphanum_characters: double (nullable = true)
|-- numberOfLinks: double (nullable = true)
|-- numwords_in_url: double (nullable = true)
|-- parametrizedLinkRatio: double (nullable = true)
|-- spelling_errors_ratio: double (nullable = true)
|-- label: double (nullable = true)
```

```
In [10]: df.select('url','alchemy_category','alchemy_category_score','is_news','label').show(10)
```

```
+-----+-----+-----+-----+-----+
|      url| alchemy_category|alchemy_category_score|is_news|label|
+-----+-----+-----+-----+-----+
|http://www.bloomb...|    business|      0.789131|    1.0| 0.0|
|http://www.popschi...|  recreation|      0.574147|    1.0| 1.0|
|http://www.menshe...|    health|      0.996526|    1.0| 1.0|
|http://www.dumbli...|    health|      0.801248|    1.0| 1.0|
|http://bleacherre...|    sports|      0.719157|    1.0| 0.0|
```

```
|http://www.conven...|      ?|      0.0| 0.0| 0.0|
|http://gofashionl...|arts_entertainment|      0.22111| 1.0| 1.0|
|http://www.inside...|      ?|      0.0| 0.0| 0.0|
|http://www.valetm...|      ?|      0.0| 1.0| 1.0|
|http://www.howswe...|      ?|      0.0| 0.0| 1.0|
+-----+-----+-----+-----+-----+
only showing top 10 rows
```

使用 randomSplit 将数据按照 7:3 的比例分成 train_df（训练数据）与 test_df（测试数据），并且.cache() 暂存在内存中，加快后续程序运行的速度。

```
In [11]: train_df,test_df = df.randomSplit([0.7,0.3])
         train_df.cache()
         test_df.cache()

Out[11]: DataFrame[url: string, alchemy_category: string, alchemy_category_score: double, avglinksize: double, com
monlinkratio_1: double, commonlinkratio_2: double, commonlinkratio_3: double, commonlinkratio_4: doub
le, compression_ratio: double, embed_ratio: double, framebased: double, frameTagRatio: double, hasDomai
nLink: double, html_ratio: double, image_ratio: double, is_news: double, lengthyLinkDomain: double, linkw
ordscore: double, news_front_page: double, non_markup_alphanum_characters: double, numberOfLinks: do
uble, numwords_in_url: double, parametrizedLinkRatio: double, spelling_errors_ratio: double, label: double]
```

1.StringIndexer

```
In [12]: from pyspark.ml.feature import StringIndexer
```

创建StringIndexer

```
In [13]: categoryIndexer = StringIndexer(
         inputCol='alchemy_category',
         outputCol='alchemy_category_Index'
         )
```

StringIndexer使用fit方法生成"Transformer"

```
In [14]: categoryTransformer=categoryIndexer.fit(df)
```

查看 categoryTransformer 的内容，categoryTransformer 的 label 属性其实就是网页分类的字 典

```
In [15]: for i in range(0,len(categoryTransformer.labels)):
         print(str(i)+'-'+categoryTransformer.labels[i])

0:?
1:recreation
2:arts_entertainment
3:business
4:health
5:sports
6:culture_politics
7:computer_internet
8:science_technology
9:gaming
10:religion
11:law_crime
```

12:unknown
13:weather

使用 `categoryTransformer` 将所有 `df` 转换为 `df1`

```
In [16]: df1 = categoryTransformer.transform(df)
```

查看转换后的 `df1` 字段

```
In [17]: df1.columns
```

```
Out[17]: ['url',  
          'alchemy_category',  
          'alchemy_category_score',  
          'avglinksiz',  
          'commonlinkratio_1',  
          'commonlinkratio_2',  
          'commonlinkratio_3',  
          'commonlinkratio_4',  
          'compression_ratio',  
          'embed_ratio',  
          'framebased',  
          'frameTagRatio',  
          'hasDomainLink',  
          'html_ratio',  
          'image_ratio',  
          'is_news',  
          'lengthyLinkDomain',  
          'linkwordscore',  
          'news_front_page',  
          'non_markup_alphanum_characters',  
          'numberOfLinks',  
          'numwords_in_url',  
          'parametrizedLinkRatio',  
          'spelling_errors_ratio',  
          'label',  
          'alchemy_category_Index']
```

查看转换后的结果

```
In [18]: df1.select('alchemy_category','alchemy_category_Index').show(10)
```

+-----+-----+	
alchemy_category	alchemy_category_Index
+-----+-----+	
business	3.0
recreation	1.0
health	4.0
health	4.0
sports	5.0
?	0.0
arts_entertainment	2.0
?	0.0
?	0.0
?	0.0
+-----+-----+	
only showing top 10 rows	

2.OneHotEncoder

导入onehotencoder模块

```
In [19]: from pyspark.ml.feature import OneHotEncoder
```

创建 OneHotEncoder

```
In [20]: encoder = OneHotEncoder(dropLast=False,
                                inputCol='alchemy_category_Index',
                                outputCol='alchemy_category_IndexVec')
```

OneHotEncoder 使用 transform 转换，结果是 df2，我们可以使用下列指令查看字段

```
In [21]: df2=encoder.transform(df1)
df2.columns
```

```
Out[21]: ['url',
'alchemy_category',
'alchemy_category_score',
'avglinksize',
'commonlinkratio_1',
'commonlinkratio_2',
'commonlinkratio_3',
'commonlinkratio_4',
'compression_ratio',
'embed_ratio',
'framebased',
'frameTagRatio',
'hasDomainLink',
'html_ratio',
'image_ratio',
'is_news',
'lengthyLinkDomain',
'linkwordscore',
'news_front_page',
'non_markup_alphanum_characters',
'numberOfLinks',
'numwords_in_url',
'parametrizedLinkRatio',
'spelling_errors_ratio',
'label',
'alchemy_category_Index',
'alchemy_category_IndexVec']
```

结果显示新增了 alchemy_category_IndexVec 字段

查看转换后新增的字段

```
In [22]: df2.select("alchemy_category","alchemy_category_Index",
                    "alchemy_category_IndexVec").show(10)

+-----+-----+-----+
| alchemy_category|alchemy_category_Index|alchemy_category_IndexVec|
+-----+-----+-----+
|      business|          3.0|    (14,[3],[1.0])|
```


recreation	1.0	(14,[1],[1.0])
health	4.0	(14,[4],[1.0])
health	4.0	(14,[4],[1.0])
sports	5.0	(14,[5],[1.0])
?	0.0	(14,[0],[1.0])
arts_entertainment	2.0	(14,[2],[1.0])
?	0.0	(14,[0],[1.0])
?	0.0	(14,[0],[1.0])
?	0.0	(14,[0],[1.0])

+-----+-----+-----+
only showing top 10 rows

3.VectorAssembler

VectorAssembler可以将多个特征字段整合成一个特征的 Vector

```
In [23]: from pyspark.ml.feature import VectorAssembler
```

创建全部特征字段 List

```
In [24]: assemblerInputs = ['alchemy_category_IndexVec']+row_df.columns[4:-1]
assemblerInputs
```

```
Out [24]: ['alchemy_category_IndexVec',
'alchemy_category_score',
'avglinksize',
'commonlinkratio_1',
'commonlinkratio_2',
'commonlinkratio_3',
'commonlinkratio_4',
'compression_ratio',
'embed_ratio',
'framebased',
'frameTagRatio',
'hasDomainLink',
'html_ratio',
'image_ratio',
'is_news',
'lengthyLinkDomain',
'linkwordscore',
'news_front_page',
'non_markup_alphanum_characters',
'numberOfLinks',
'numwords_in_url',
'parametrizedLinkRatio',
'spelling_errors_ratio']
```

创建 VectorAssembler

```
In [25]: assembler = VectorAssembler(inputCols=assemblerInputs,
outputCol="features")
```

运行 VectorAssembler 转换

```
In [26]: df3=assembler.transform(df2)
```

查看整合后的新增字段

```
In [27]: df3.columns

Out [27]: ['url',
            'alchemy_category',
            'alchemy_category_score',
            'avglinksiz',
            'commonlinkratio_1',
            'commonlinkratio_2',
            'commonlinkratio_3',
            'commonlinkratio_4',
            'compression_ratio',
            'embed_ratio',
            'framebased',
            'frameTagRatio',
            'hasDomainLink',
            'html_ratio',
            'image_ratio',
            'is_news',
            'lengthyLinkDomain',
            'linkwordscore',
            'news_front_page',
            'non_markup_alphanum_characters',
            'numberOfLinks',
            'numwords_in_url',
            'parametrizedLinkRatio',
            'spelling_errors_ratio',
            'label',
            'alchemy_category_Index',
            'alchemy_category_IndexVec',
            'features']
```

```
In [28]: df3.select('features').take(1)

Out [28]: [Row(features=SparseVector(36, {3: 1.0, 14: 0.7891, 15: 2.0556, 16: 0.6765, 17: 0.2059, 18: 0.0471, 19: 0.0
235, 20: 0.4438, 23: 0.0908, 25: 0.2458, 26: 0.0039, 27: 1.0, 28: 1.0, 29: 24.0, 31: 5424.0, 32: 170.0, 33: 8.0,
34: 0.1529, 35: 0.0791}))]
```

4.DecisionTreeClassifier

```
In [29]: from pyspark.ml.classification import DecisionTreeClassifier
```

运行 DesionTreeClassifier

```
In [30]: dt = DecisionTreeClassifier(labelCol="label",featuresCol="features",
                                     impurity="gini",maxDepth=10,maxBins=14)
```

进行训练，之前创建的 dt 决策树分类，我们可以使用.fit() 方法进行训练，训练结果产生 dt_model 模型，之后可以使用 print 查看产生的模型

```
In [31]: dt_model=dt.fit(df3)
         print(dt_model)
```

DecisionTreeClassificationModel (uid=DecisionTreeClassifier_2a7f3eec7817) of depth 10 with 653 nodes

进行预测，建立模型后就可以使用`.trainform()` 进行转换了，转换后会产生预测结果 `df4/`

```
In [32]: df4=dt_model.transform(df3)
```

5 建立机器学习 Pipeline 流程

```
In [33]: from pyspark.ml import Pipeline
from pyspark.ml.feature import StringIndexer, OneHotEncoder, VectorAssembler
from pyspark.ml.classification import DecisionTreeClassifier
```

建立 pipeline

```
In [34]: stringIndexer = StringIndexer(inputCol='alchemy_category',
outputCol="alchemy_category_Index")
encoder = OneHotEncoder(dropLast=False,
inputCol='alchemy_category_Index',
outputCol="alchemy_category_IndexVec")
assemblerInputs = ['alchemy_category_IndexVec'] + row_df.columns[4:-1]
assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
dt = DecisionTreeClassifier(labelCol="label", featuresCol="features", impurity="gini",
maxDepth=10, maxBins=14)

pipeline = Pipeline(stages=[stringIndexer, encoder, assembler, dt])
```

```
In [35]: pipeline.getStages()
```

```
Out [35]: [StringIndexer_bd31abc032f4,
OneHotEncoder_97883791e25c,
VectorAssembler_96032e56ba91,
DecisionTreeClassifier_07bfaed33184]
```

6.使用 pipeline 进行数据处理与训练

```
In [36]: pipelineModel = pipeline.fit(train_df)
```

看训练完成后的决策树模型

```
In [37]: pipelineModel.stages[3]
```

```
Out [37]: DecisionTreeClassificationModel (uid=DecisionTreeClassifier_07bfaed33184) of depth 10 with 615 nodes
```

查看训练完后的决策树模型规则

```
In [38]: print (pipelineModel.stages[3].toDebugString[:1000])
```

```
DecisionTreeClassificationModel (uid=DecisionTreeClassifier_07bfaed33184) of depth 10 with 615 nodes
If (feature 31 <= 1765.0)
  If (feature 2 in {1.0})
    If (feature 26 <= 0.201427045)
      If (feature 20 <= 0.484175714)
        If (feature 29 <= 21.5)
          Predict: 0.0
```

```
Else (feature 29 > 21.5)
  If (feature 34 <= 0.049220706)
    If (feature 15 <= 3.723642426)
      If (feature 15 <= 1.0132735585)
        Predict: 1.0
      Else (feature 15 > 1.0132735585)
        If (feature 20 <= 0.4488383145)
          Predict: 0.0
        Else (feature 20 > 0.4488383145)
          Predict: 1.0
      Else (feature 15 > 3.723642426)
        Predict: 1.0
    Else (feature 34 > 0.049220706)
      If (feature 14 <= 0.8503704999999999)
        If (feature 29 <= 65.5)
          Predict: 1.0
        Else (feature 29 > 65.5)
          Predict: 0.0
      Else (feature 14 > 0.8503704999999999)
        Predict: 0.0
    Else (feature 20 > 0.484175714)
      If (
```

7.使用 pipelineModel 进行预测

使用 pipelineModel 的 transform 方法，传入 test_df 测试数据进行预测

```
In [39]: predicted=pipelineModel.transform(test_df)
```

查看预测后的 Schema，发现新增了 3 个字段

```
In [40]: predicted.columns
```

```
Out[40]: ['url',
'alchemy_category',
'alchemy_category_score',
'avglinksiz',
'commonlinkratio_1',
'commonlinkratio_2',
'commonlinkratio_3',
'commonlinkratio_4',
'compression_ratio',
'embed_ratio',
'framebased',
'frameTagRatio',
'hasDomainLink',
'html_ratio',
'image_ratio',
'is_news',
'lengthyLinkDomain',
'linkwordscore',
'news_front_page',
'non_markup_alphanum_characters',
'numberOfLinks',
'numwords_in_url',
'parametrizedLinkRatio',
'spelling_errors_ratio',
'label',
```

```
'alchemy_category_Index',
'alchemy_category_IndexVec',
'features',
'rawPrediction',
'probability',
'prediction']
```

看预测结果 DataFrame

```
In [41]: predicted.select('url','features','rawprediction','probability','label','prediction').show(10)
```

url	features	rawprediction	probability	label	prediction
http://1000awesom...	(36,[0,15,16,17,1...	[34.0,163.0]	[0.17258883248730...	1.0	1.0
http://100miledie...	(36,[0,15,20,23,2...	[1.0,13.0]	[0.07142857142857...	0.0	1.0
http://13gb.com/v...	(36,[0,15,16,20,2...	[56.0,30.0]	[0.65116279069767...	0.0	0.0
http://3kidsandus...	(36,[0,15,16,17,1...	[34.0,163.0]	[0.17258883248730...	0.0	1.0
http://3kidsandus...	(36,[0,15,16,17,1...	[34.0,163.0]	[0.17258883248730...	1.0	1.0
http://3kidsandus...	(36,[0,15,16,17,1...	[55.0,29.0]	[0.65476190476190...	0.0	0.0
http://8tracks.co...	(36,[2,14,15,16,1...	[0.0,3.0]	[0.0,1.0]	0.0	1.0
http://9gag.com/g...	(36,[0,15,16,17,1...	[2.0,11.0]	[0.15384615384615...	0.0	1.0
http://abeautiful...	(36,[2,14,15,16,1...	[29.0,3.0]	[0.90625,0.09375]	1.0	0.0
http://addapinch....	(36,[1,14,15,16,1...	[3.0,105.0]	[0.02777777777777...	0.0	1.0

only showing top 10 rows

查看预测结果与概率

```
In [42]: predicted.select('probability','prediction') .take(10)
```

```
Out [42]: [Row(probability=DenseVector([0.1726, 0.8274]), prediction=1.0),
Row(probability=DenseVector([0.0714, 0.9286]), prediction=1.0),
Row(probability=DenseVector([0.6512, 0.3488]), prediction=0.0),
Row(probability=DenseVector([0.1726, 0.8274]), prediction=1.0),
Row(probability=DenseVector([0.1726, 0.8274]), prediction=1.0),
Row(probability=DenseVector([0.6548, 0.3452]), prediction=0.0),
Row(probability=DenseVector([0.0, 1.0]), prediction=1.0),
Row(probability=DenseVector([0.1538, 0.8462]), prediction=1.0),
Row(probability=DenseVector([0.9062, 0.0938]), prediction=0.0),
Row(probability=DenseVector([0.0278, 0.9722]), prediction=1.0)]
```

8. 评估模型的准确率

首先从 `pyspark.ml.evaluation` 导入 `BinaryClassificationEvaluator` 模块

```
In [43]: from pyspark.ml.evaluation import BinaryClassificationEvaluator
```

创建 `BinaryClassificationEvaluator`，传入下列参数： - `rawPredictionCol`="rawPrediction" 之前预测后产生的字段

- `labelCol`="label" 标签字段
- `metricName`="areaUnderROC" 也就是 AUC

```
In [44]: evaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction",
                                                labelCol="label",
                                                metricName="areaUnderROC" )
```

计算AUC

```
In [45]: predictions = pipelineModel.transform(test_df)
auc= evaluator.evaluate(predictions)
auc
```

```
Out [45]: 0.6324079757381753
```

使用 `TrainValidation` 进行训练验证找出最佳模型

从 `pyspark.ml.tuning` 导入 `ParamGridBuilder` 与 `TrainValidationSplit` 模块

```
In [46]: from pyspark.ml.tuning import ParamGridBuilder,TrainValidationSplit
```

设置训练验证的参数，我们使用 `ParamGridBuilder` 设置 `impurity` 两个参数值、`maxDepth` 三个参数值与 `maxBins` 三个参数值，后续执行训练验证时会执行 **233=18** 次。

```
In [47]: paramGrid = ParamGridBuilder()\
.addGrid(dt.impurity, [ "gini","entropy"])\
.addGrid(dt.maxDepth, [ 5,10,15])\
.addGrid(dt.maxBins, [10, 15,20])\
.build()
```

创建 `TrainValidationSplit`，传入下列参数，执行后创建 `tv` 变量：

- `estimator=dt`，之前创建的 `DecisionTreeClassifier`
- `evaluator=evaluator`，之前创建的 `BinaryClassificationEvaluator`
- `estimatorParamMaps=paramGrid`，之前创建的 `ParamGridBuilder`
- `trainRatio=0.8`，训练验证前会先将数据按照 **8:2** 的比例分成训练数据与验证数据

```
In [48]: tvs = TrainValidationSplit(estimator=dt,evaluator=evaluator,
                                   estimatorParamMaps=paramGrid,trainRatio=0.8)
```

建立 `tvs_pipeline/`

```
In [49]: tvs_pipeline = Pipeline(stages=[stringIndexer,encoder ,assembler, tvs])
```

使用 `tvs_pipeline` 流程进行训练验证

```
In [50]: tvs_pipelineModel =tvs_pipeline.fit(train_df)
```

查看训练完成的最佳模型

```
In [51]: bestModel=tvs_pipelineModel.stages[3].bestModel
bestModel
```

```
Out [51]:
```

DecisionTreeClassificationModel (uid=DecisionTreeClassifier_07bfaed33184) of depth 15 with 1255 nodes

看训练验证完成的最佳模型规则, [: 500] 表示只显示前 500 文字

```
In [52]: print (bestModel.toDebugString[:500])

DecisionTreeClassificationModel (uid=DecisionTreeClassifier_07bfaed33184) of depth 15 with 1255 nodes
If (feature 31 <= 1548.0)
  If (feature 2 in {1.0})
    If (feature 33 <= 11.5)
      If (feature 33 <= 5.5)
        If (feature 35 <= 0.128495221)
          If (feature 28 <= 0.5)
            If (feature 16 <= 0.5052088985000001)
              If (feature 17 <= 0.103194726)
                If (feature 26 <= 0.6825757575)
                  If (feature 34 <= 0.179858142)
                    If (feature 29 <= 22.5)
                      I
```

评估最佳模型 AUC

```
In [53]: predictions = tvs_pipelineModel.transform(test_df)
auc= evaluator.evaluate(predictions)
auc

Out[53]: 0.6440101246327793
```

9.使用 **crossValidation** 交叉验证找出最佳模型

```
In [54]: from pyspark.ml.tuning import CrossValidator
```

建立交叉验证的 **CrossValidator**

```
In [55]: cv = CrossValidator(estimator=dt, evaluator=evaluator,
                           estimatorParamMaps=paramGrid, numFolds=3)
```

建立交叉验证的 **cv_pipeline**

```
In [56]: cv_pipeline = Pipeline(stages=[stringIndexer,encoder ,assembler, cv])
```

用 **cv_pipeline** 流程进行交叉验证

```
In [57]: cv_pipelineModel = cv_pipeline.fit(train_df)
```

查看交叉验证完成的最佳模型

```
In [58]: bestModel=cv_pipelineModel.stages[3].bestModel
bestModel

Out[58]: DecisionTreeClassificationModel (uid=DecisionTreeClassifier_07bfaed33184) of depth 15 with 1575 nodes
```

评估最佳模型 AUC

```
In [59]: predictions = cv_pipelineModel.transform(test_df)
auc= evaluator.evaluate(predictions)
auc
```

Out [59]: 0.6608278234025085

使用随机森林 RandomForestClassifier 分类器

创建 RandomForestClassifier 变量 rf，传入参数与决策树类似，只是多了 numTrees 参数（设置 决策森林中有多少决策树，这里设为 10）

```
In [60]: from pyspark.ml.classification import RandomForestClassifier
rf=RandomForestClassifier(labelCol="label",featuresCol="features",numTrees=10)
rfpipeline = Pipeline(stages=[stringIndexer,encoder ,assembler,rf ])
```

评估 RandomForestClassifier 的准确度

rfpipeline.fit 传入 train_df 进行训练，再用 rftvs_pipelineModel.transform 传入 test_df 进行评估，我们可以看到 AUC 约为 0.72，比之前使用决策树的准确度明显增加

```
In [61]: rfpipelineModel = rfpipeline.fit(train_df)
rfpredicted=rfpipelineModel.transform(test_df)
evaluator.evaluate(rfpredicted)
```

Out [61]: 0.7371290346808376

使用 RandomForestClassifier TrainValidation 找出最佳模型

```
In [62]: from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.classification import RandomForestClassifier
paramGrid = ParamGridBuilder()\
.addGrid(rf.impurity, [ "gini","entropy"])\
.addGrid(rf.maxDepth, [ 5,10,15])\
.addGrid(rf.maxBins, [10, 15,20])\
.addGrid(rf.numTrees, [10, 20,30])\
.build()
rftvs = TrainValidationSplit(estimator=rf, evaluator=evaluator,
estimatorParamMaps=paramGrid, trainRatio=0.8)
rftvs_pipeline = Pipeline(stages=[stringIndexer,encoder ,assembler, rftvs])
rftvs_pipelineModel =rftvs_pipeline.fit(train_df)
rftvspredictions = rftvs_pipelineModel.transform(test_df)
auc= evaluator.evaluate(rftvspredictions)
auc
```

Out [62]: 0.7653783831147928

使用 crossValidation 找出最佳模型

```
In [63]: from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
from pyspark.ml import Pipeline
rfcv = CrossValidator(estimator=rf, evaluator=evaluator,
estimatorParamMaps=paramGrid, numFolds=3)
```



```
rfcv_pipeline = Pipeline(stages=[stringIndexer,encoder ,assembler, rfcv])
rfcv_pipelineModel = rfcv_pipeline.fit(train_df)
```

使用最佳模型进行预测

```
In [64]: rfcvpredictions = rfcv_pipelineModel.transform(test_df)
```

显示使用最佳模型进行预测结果

```
In [67]: DescDict = {
0: "暂时性网页 (ephemeral)",
1: "长青网页 (evergreen)"
}
for data in rfcvpredictions.select('url','prediction').take(5):
    print(" 网址: " +str(data[0])+"\n" +\
          " ==> 预测:" + str(data[1])+\
          " 说明:" +DescDict[data[1]] +"\n")
```

网址: <http://1000awesomethings.com/2008/07/07/989-blowing-your-nose-in-the-shower/>
 ==> 预测:1.0 说明:长青网页 (evergreen)

网址: <http://100milediet.org/>
 ==> 预测:0.0 说明:暂时性网页 (ephemeral)

网址: <http://13gb.com/videos/2298/>
 ==> 预测:0.0 说明:暂时性网页 (ephemeral)

网址: <http://3kidsandus.com/2010/patriotic-checkerboard-cake-for-the-4th-of-july/>
 ==> 预测:1.0 说明:长青网页 (evergreen)

网址: <http://3kidsandus.com/2011/cook-for-the-cure-strawberries-n-cream-layer-cake/>
 ==> 预测:1.0 说明:长青网页 (evergreen)

计算最佳模型 AUC

```
In [68]: auc= evaluator.evaluate(rfcvpredictions)
auc
```

```
Out [68]: 0.7632262080194218
```